

In `simplebuffer.c`, the vulnerability stems from the size of the char array `buff`, and the user input. In the function `helper()`, the program only allocates space for 4 char elements, and correspondingly 4 bytes in memory for the `buff` array as a char takes up 1 byte of space. However, in the `fgets` call, the program prompts the user for an input that can be up to `INT_MAX` elements (which is much greater than 4) and stores each character of the input in `buff`. This means that there is no bound check/stop for the input, and an input greater than 4 characters would write into other parts of memory. In exploiting the program, it required understanding how the stack for an x86 machine is laid out and then crafting the input to `simplebuffer` such that the desired code is run. In x86, the return address of the current function is stored right before the current call frame's frame pointer (RBP), and the arguments for the current function are stored right before the return address. So, when I pass in the shellcode as input to the `helper`, I need to find the exact memory address of where the spawn code would be stored, and then overwrite the return address with this memory address. This required some “padding” input to overwrite the stack until the return address, then calculating where exactly the shellcode begins and putting that memory address as the return address of the `helper` function.

```
# x86 shellcode
addr= '\xb0\x56\xff\xff\xff\x7f\x00\x00'
prepend= '\x00\x00\x00\x00\x00'
shellcode= '\x48\x83\xec\x20\x48\x89\xe5\x6a' \
           '\x3b\x58\x48\x31\xff\x57\x48\xb8' \
           '\x2f\x62\x69\x6e\x2f\x2f\x73\x68' \
           '\x57\x48\x89\xe7\x48\x31\xf6\x56' \
           '\x57\x48\x89\xe6\x48\x31\xd2\x0f\x05';

# TODO actually generate the other parts of the input necessary to
# execute the buffer overflow, combine it with the shellcode, and print it all out

print prepend+prepend+prepend+prepend+addr+shellcode
```

I determined that there were 20 bytes between the address of `buff[0]` and the return address, so I created a padding/constant string called “prepend” and prepended the final string with 20 bytes of “junk”. Now the next 8 bytes would overwrite the return address (i.e. where the function returns to), and since I wanted the function to return the beginning of where the spawn code starts (in the stack this is immediately before the return address). I discovered that the spawn code starts at address `0x7fffffff56b0`, so I overwrote the return address with this address.

